



Chapter 9: Deep Neural Networks

EE2405

嵌入式系統與實驗


Embedded System Lab





Machine Learning Overview

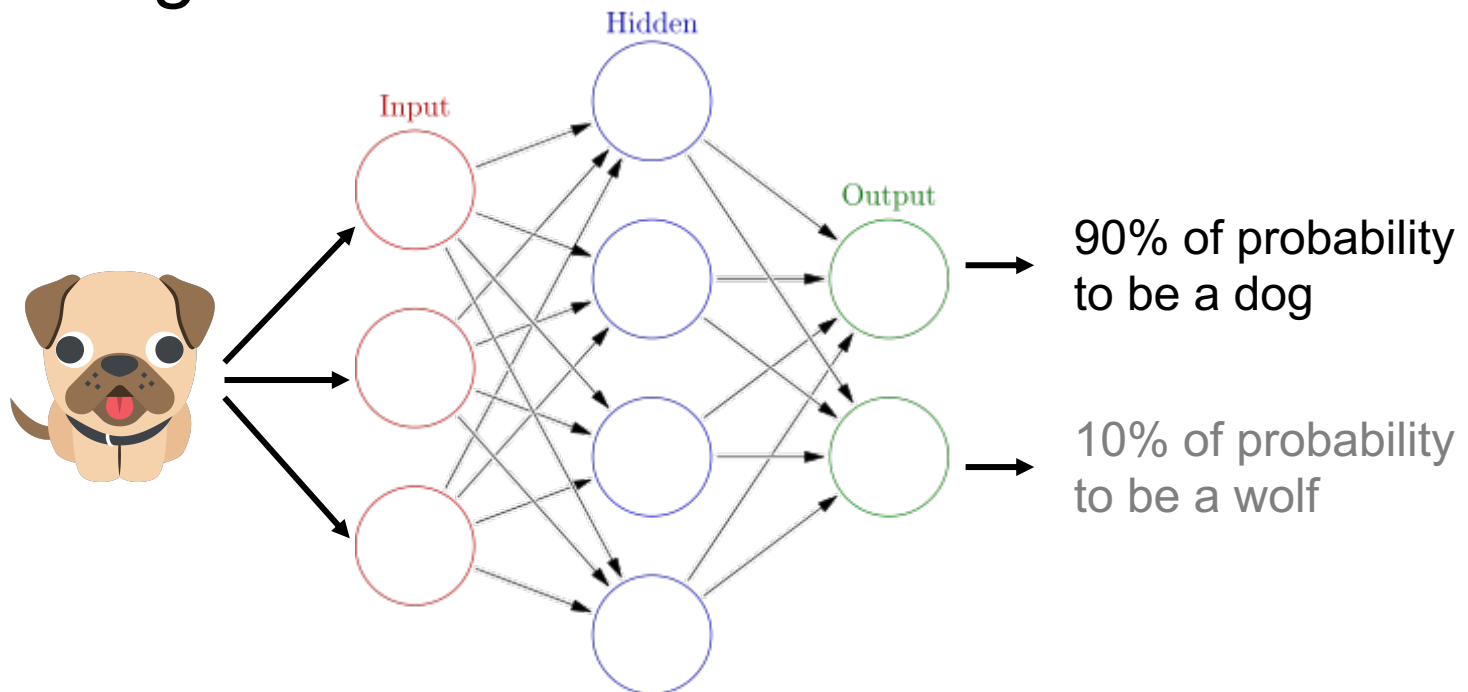
- Supervised: Learning with a labeled training set
 - Example: classify animal images
- Unsupervised: Discover patterns in unlabeled data
 - Example: cluster people with similar buying habits





Deep Learning


- One of machine learning techniques
- Apply multi-layer neural networks to machine learning





ImageNet Competition


- Convolutional neural networks improve image classification dramatically





Neuron Model

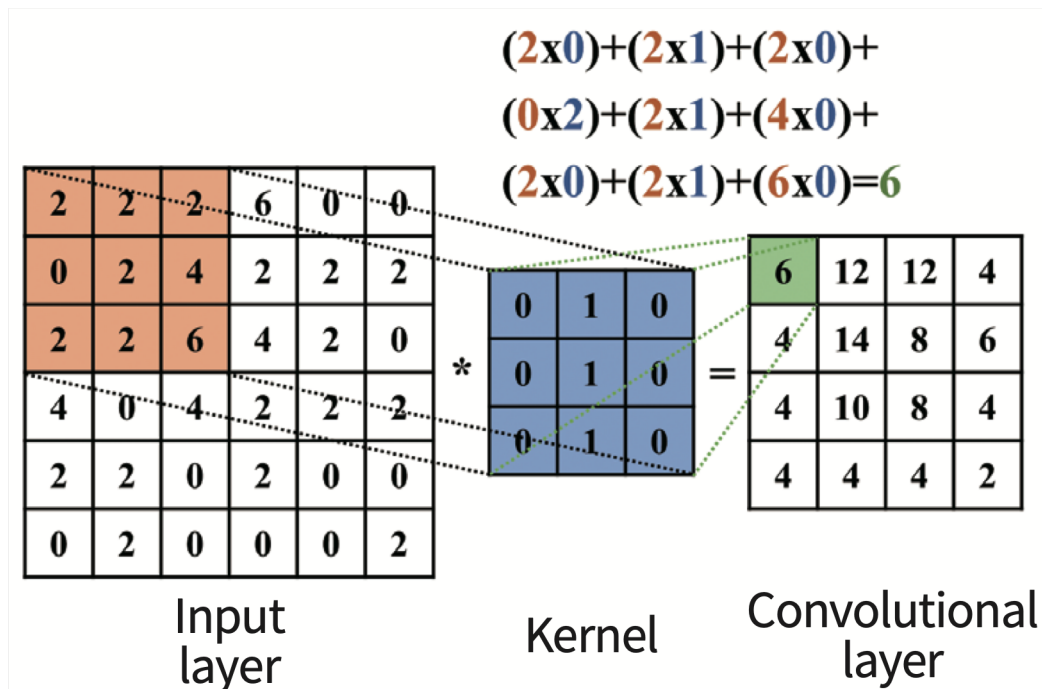
- A neuron computes weighted sum with bias and pass through an activation function to output





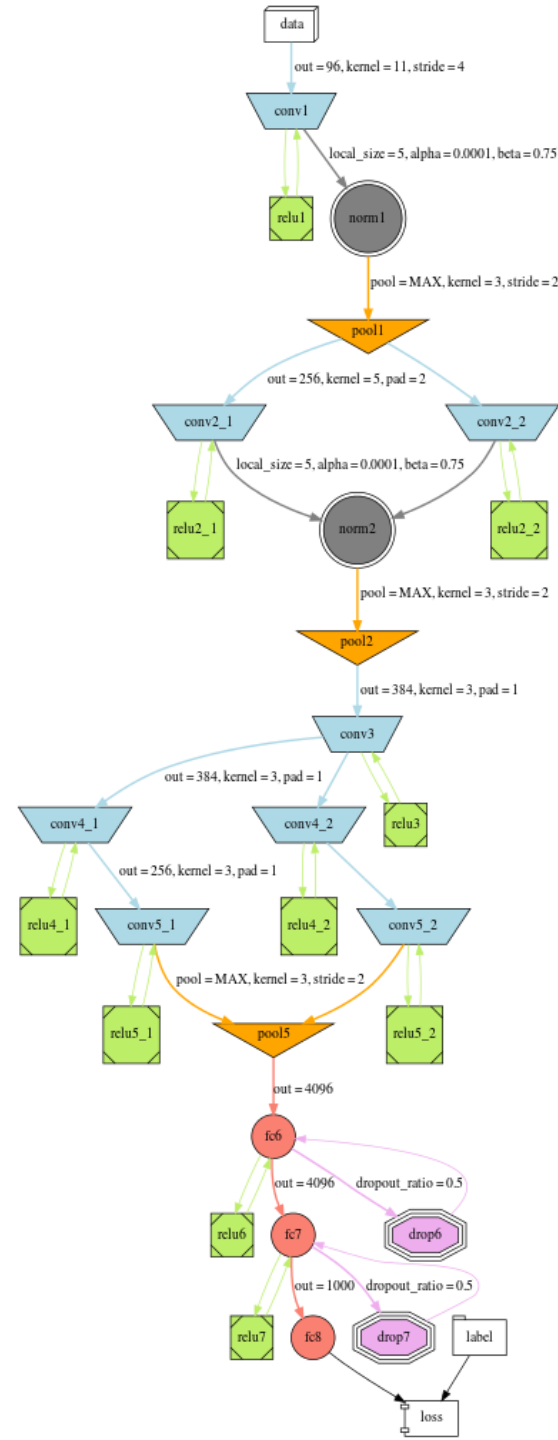
Convolutional Neural Networks

- CNN computes multi-dimension convolution instead of 1-D weighted sums
- For image processing 2-D filters are common



AlexNet

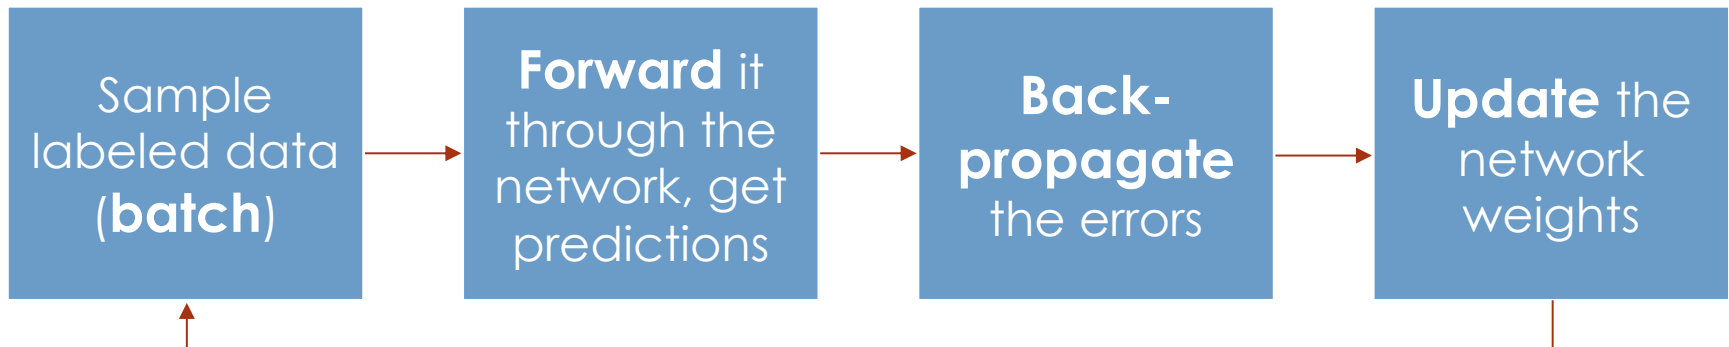
- Convolution
- Relu
- Normalization
- Pooling
- Fully-connected (FC)






Training a Network

- The goal is to optimize (reduce) output loss
 - Need several rounds (epoch)
- Output loss is the difference between known output (100% a dog) and actual prediction (90% a dog)



Repeat till convergence



Discussion on CNN Training

- Prepare labeled (tagged) data
 - For both training and testing
- Use mini-batch
- Avoid overfitting (use dropouts)
- Hyperparameters




Training Frameworks

- In the beginning, Numpy...
- Need more computation resources.
Frameworks start to code with GPU: Cafe, Tensorflow, MxNet etc.
 - For specialized networks: Yolo
- To reduce programming efforts: Keras, PyTorch, Tensorflow 2.0, etc.
 - All with Python interface.



A Training Example

- To train a network to model data generated by a sine function.
 - Data from a sine + some noise with a fix amp/freq
- The example source:
 - https://colab.research.google.com/github/tensorflow/tensorflow/blob/master/tensorflow/lite/micro/examples/hello_world/create_sine_model.ipynb



Training Steps

- Generate data with noise
- Split data into [training, validation, testing]
- Design a model
- Train the model
- Revise the model
- Retrain



1. Generate data with noise

```
SAMPLES = 1000
```

```
np.random.seed(1337)
```

```
# Generate a uniformly distributed set of random numbers in [0 to 2π)  
x_values = np.random.uniform(low=0, high=2*math.pi, size=SAMPLES)
```

```
# Shuffle the values to guarantee they're not in order  
np.random.shuffle(x_values)
```

```
# Calculate the corresponding sine values
```

```
y_values = np.sin(x_values)
```

```
# *operator unpack a list as individual parameters
```



```
# randn() to generate data of std normal distribution in y_values.shape
```

```
y_values += 0.1 * np.random.randn(*y_values.shape)
```

```
# The 'b.' argument tells the library to print blue dots.
```

```
plt.plot(x_values, y_values, 'b.')
```

```
plt.show()
```







Split Data

```
# 60% for training and 20% for testing, 20% for validation.
TRAIN_SPLIT = int(0.6 * SAMPLES)
TEST_SPLIT = int(0.2 * SAMPLES + TRAIN_SPLIT)

# Use np.split to chop our data into three parts.
# Two indices in 2nd arg, so the data will be divided into three chunks.
x_train, x_test, x_validate = np.split(x_values, [TRAIN_SPLIT,
TEST_SPLIT])
y_train, y_test, y_validate = np.split(y_values, [TRAIN_SPLIT,
TEST_SPLIT])

# Double check that our splits add up correctly
assert (x_train.size + x_validate.size + x_test.size) == SAMPLES

# Plot the data in each partition in different colors:
plt.plot(x_train, y_train, 'b.', label="Train")
plt.plot(x_test, y_test, 'r.', label="Test")
plt.plot(x_validate, y_validate, 'y.', label="Validate")
plt.legend()
plt.show()
```





Design a Model

```
# TF 2.0 Keras to create a simple model architecture  
# Sequential() is a layer by layer model  
from tensorflow.keras import layers  
model_1 = tf.keras.Sequential()  
  
# First layer takes a scalar input and feeds it through 16 "neurons". The  
# neurons decide whether to activate based on the 'relu' activation  
function.  
model_1.add(layers.Dense(16, activation='relu', input_shape=(1,)))  
  
# Final layer is a single neuron, since we want to output a single value  
model_1.add(layers.Dense(1))  
  
# Compile the model using a standard optimizer and loss function for  
regression  
model_1.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
```



Train the Model

Train the model on our training data while validating on our validation set

```
history_1 = model_1.fit(x_train, y_train,  
epochs=1000, batch_size=16,  
validation_data=(x_validate, y_validate))
```

Train on 600 samples, validate on 200 samples

Epoch 1/1000

600/600 [=====] - 0s 412us/sample - loss: 0.5016 - mae: 0.6297 - val_loss: 0.4922 - val_mae: 0.6235

Epoch 2/1000

600/600 [=====] - 0s 105us/sample - loss: 0.3905 - mae: 0.5436 - val_loss: 0.4262 - val_mae: 0.5641

...

Epoch 998/1000

600/600 [=====] - 0s 109us/sample - loss: 0.1535 - mae: 0.3068 - val_loss: 0.1507 - val_mae: 0.3113

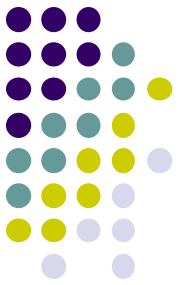
Epoch 999/1000

600/600 [=====] - 0s 100us/sample - loss: 0.1545 - mae: 0.3077 - val_loss: 0.1499 - val_mae: 0.3103

Epoch 1000/1000

600/600 [=====] - 0s 132us/sample - loss: 0.1530 - mae: 0.3045 - val_loss: 0.1542 - val_mae: 0.3143

Plot Training and Validation Loss




*# Draw a graph of the loss, which is the distance between
the predicted and actual values during training and
validation.*

```
loss = history_1.history['loss']  
val_loss = history_1.history['val_loss']  
  
epochs = range(1, len(loss) + 1)  
  
plt.plot(epochs, loss, 'g.', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



Training and validation loss





Skip 50 Epocs


Exclude the first few epochs so the graph is easier to read

SKIP = 50

```
plt.plot(epochs[SKIP:], loss[SKIP:], 'g.',  
label='Training loss')  
plt.plot(epochs[SKIP:], val_loss[SKIP:], 'b.',  
label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



Training and validation loss






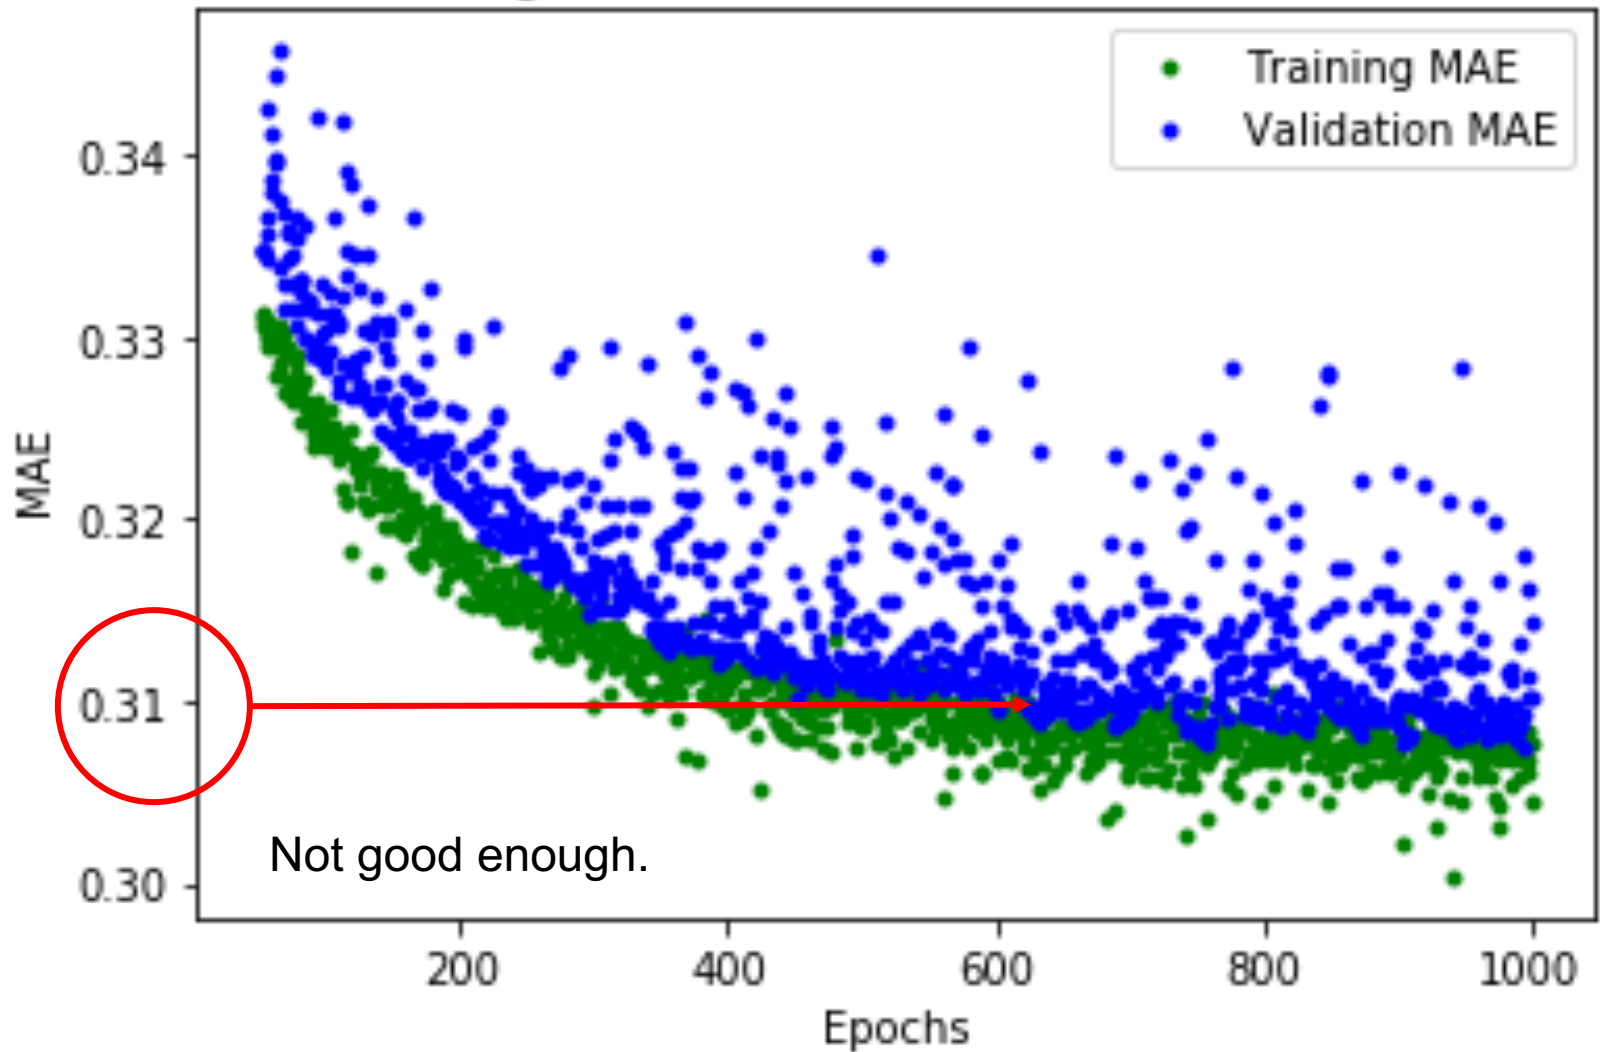
Plot Mean Average Error

```
# Draw a graph of mean absolute error, which is another
way of
# measuring the amount of error in the prediction.
mae = history_1.history['mae']
val_mae = history_1.history['val_mae']

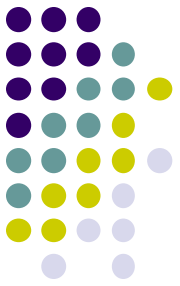
plt.plot(epochs[SKIP:], mae[SKIP:], 'g.', label='Training
MAE')
plt.plot(epochs[SKIP:], val_mae[SKIP:], 'b.',
label='Validation MAE')
plt.title('Training and validation mean absolute error')
plt.xlabel('Epochs')
plt.ylabel('MAE')
plt.legend()
plt.show()
```



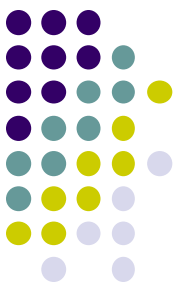
Training and validation mean absolute error




Print Predicted Values (Inference on Training Data)



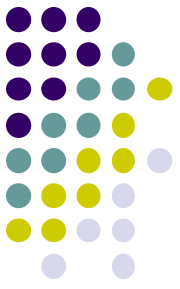
```
# Use the model to make predictions from our  
validation data  
predictions = model_1.predict(x_train)  
  
# Plot the predictions along with to the test data  
plt.clf()  
plt.title('Training data predicted vs actual  
values')  
plt.plot(x_test, y_test, 'b.', label='Actual')  
plt.plot(x_train, predictions, 'r.',  
label='Predicted')  
plt.legend()  
plt.show()
```



Training data predicted vs actual values



Add another Layer to the Model



```
model_2 = tf.keras.Sequential()
```

```
# First layer takes a scalar input and feeds it through 16 "neurons". The  
# neurons decide whether to activate based on the 'relu' activation  
function.
```

```
model_2.add(layers.Dense(16, activation='relu', input_shape=(1,)))
```

```
# The new second layer may help the network learn more complex  
representations
```

```
model_2.add(layers.Dense(16, activation='relu'))
```

```
# Final layer is a single neuron, since we want to output a single value
```


```
model_2.add(layers.Dense(1))
```


```
# Compile the model using a standard optimizer and loss function for  
regression
```

```
model_2.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
```




Training and validation mean absolute error





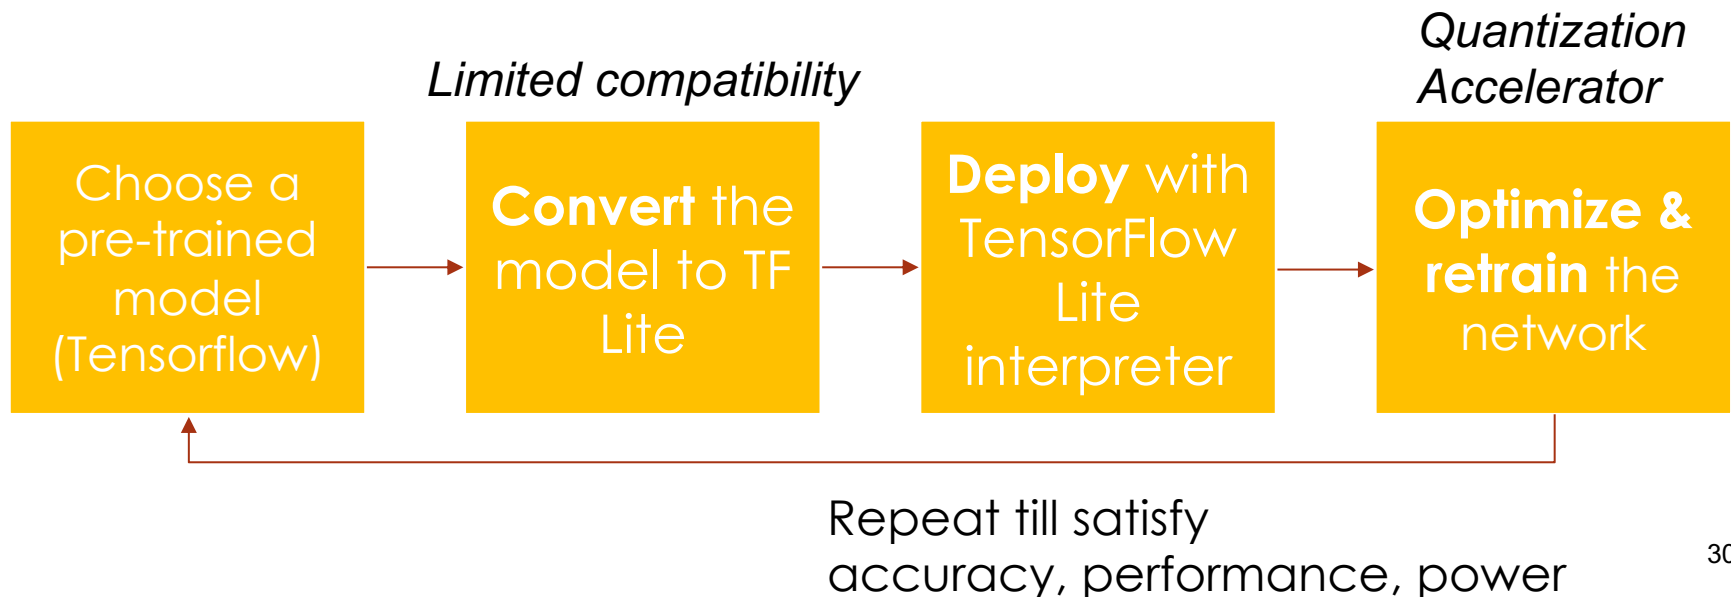
Comparison of predictions and actual values



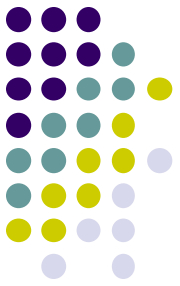


Tensorflow Lite

- TensorFlow Lite is an open source deep learning framework for on-device **inference**.
 - Android and IOS devices
 - Micro-controllers

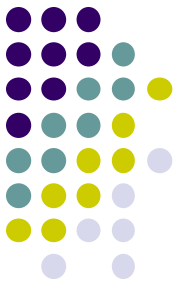


Tensorflow Lite on Micro-controller



- https://www.tensorflow.org/lite/microcontrollers/get_started

Convert to TensorFlow Lite Model



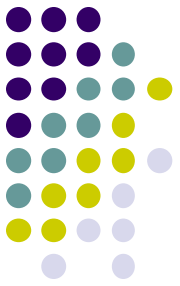
```
# Convert the model to the TensorFlow Lite format  
without quantization
```

```
converter =  
tf.lite.TFLiteConverter.from_keras_model(model_2)  
tflite_model = converter.convert()
```

```
# Save the model to disk
```

```
open("sine_model.tflite", "wb").write(tflite_model)
```


Convert to TensorFlow Lite Quantized Model



```
# Convert the model to the TensorFlow Lite format with quantization
converter = tf.lite.TFLiteConverter.from_keras_model(model_2)
converter.optimizations =
[tf.lite.Optimize.OPTIMIZE_FOR_SIZE]
tflite_model = converter.convert()

# Save the model to disk
open("sine_model_quantized.tflite", "wb").write(tflite_model)
```



Inference in Tensorflow Lite

```
# Instantiate an interpreter for each model
sine_model = tf.lite.Interpreter('sine_model.tflite')

# Allocate memory for each model
sine_model.allocate_tensors()


# Get the input and output tensors so we can feed in values and get the results
sine_model_input =
sine_model.tensor(sine_model.get_input_details()[0]["index"])
sine_model_output =
sine_model.tensor(sine_model.get_output_details()[0]["index"])

# Create arrays to store the results
sine_model_predictions = np.empty(x_test.size)

# Run each model's interpreter for each value and store the results in arrays
for i in range(x_test.size):
    sine_model_input().fill(x_test[i])
    sine_model.invoke()
    sine_model_predictions[i] = sine_model_output()[0]
```



Comparison of various models against actual values





Write to a C File

- Install xxd

```
$ apt-get -qq install xxd
```

- Use xxd to convert tflite file to a C hex file

```
$ xxd -i sine_model_quantized.tflite >  
sine_model_quantized.cc
```

```
unsigned char sine_model_quantized_tflite[] = {  
0x18, 0x00, 0x00, 0x00, 0x54, 0x46, 0x4c, 0x33, 0x00, 0x00, 0x0e, 0x00,  
0x18, 0x00, 0x04, 0x00, 0x08, 0x00, 0x0c, 0x00, 0x10, 0x00, 0x14, 0x00,  
...
```

Inference on Micro-controller

(1) Includes



```
// sine_model_data.h generated from TF Lite model  
// all_ops_resolver.h provides the operations used by the interpreter to  
run the model.  
// micro_error_reporter.h outputs debug information.  
// micro_interpreter.h contains code to load and run models.  
// schema_generated.h contains the schema for the TensorFlow Lite  
FlatBuffer model file format.  
// version.h provides versioning information for the TensorFlow Lite  
schema.
```

```
#include "tensorflow/lite/micro/examples/hello_world/sine_model_data.h"  
#include "tensorflow/lite/micro/kernels/all_ops_resolver.h"  
#include "tensorflow/lite/micro/micro_error_reporter.h"  
#include "tensorflow/lite/micro/micro_interpreter.h"  
#include "tensorflow/lite/micro/testing/micro_test.h"  
#include "tensorflow/lite/schema/schema_generated.h"  
#include "tensorflow/lite/version.h"
```



Set up Logging

- Setup the error/output mechanism for each micro-controller.

```
tflite::MicroErrorReporter micro_error_reporter;  
tflite::ErrorReporter* error_reporter =  
&micro_error_reporter;
```



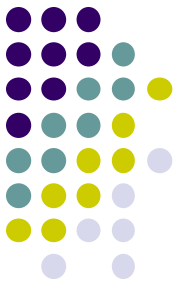
Load the Model

```
// Map the model into a usable data structure.  
// This doesn't involve any copying or parsing,  
// it's a very lightweight operation.
```

```
const tflite::Model* model =  
::tflite::GetModel(g_sine_model_data);
```

```
if (model->version() != TFLITE_SCHEMA_VERSION) {  
    TF_LITE_REPORT_ERROR(error_reporter,  
        "Model provided is schema version %d not equal "  
        "to supported version %d.\n",  
        model->version(), TFLITE_SCHEMA_VERSION);  
}
```

Instantiate Operations Resolver



- To call operators in TF Lite, we need to create a resolver object.

```
// This pulls in all the operation implementations  
in TF Lite
```

```
tflite::ops::micro::AllOpsResolver resolver;
```

- To save memory, and only pull in necessary operators, use `MicroMutableOpResolver`



Allocate Tensor Memories

```
// Create an area of memory to use for  
input, output, and intermediate arrays.  
// Finding the minimum value for your  
model may require some trial and error.  
const int tensor_arena_size = 2 * 1024;  
uint8_t tensor_arena[tensor_arena_size];
```



Instantiate Interpreter

```
// Build an interpreter to run the model  
tflite::MicroInterpreter interpreter(model,  
resolver, tensor_arena, tensor_arena_size,  
error_reporter);
```



Allocate Tensors

- Ask the interpreter to allocate memory from the tensor_arena for the model's tensors

```
interpreter.AllocateTensors();
```

- The MicroInterpreter instance can provide us with a pointer to the model's input tensor by calling `.input(0)`, where 0 represents the first (and only) input tensor:

```
// Obtain a pointer to the model's input tensor  
TfLiteTensor* input = interpreter.input(0);
```



Validate Input Shape

- Use TF Lite unit test framework

```
// Make sure the input has the properties we expect
TF_LITE_MICRO_EXPECT_NE(nullptr, input);
// The property "dims" tells us the tensor's shape. It has one element for
// each dimension. Our input is a 2D tensor containing 1 element,
// so "dims" should have size 2.
TF_LITE_MICRO_EXPECT_EQ(2, input->dims->size);
// The value of each element gives the length of the corresponding tensor.
// We should expect two single element tensors
// (one is contained within the other).
TF_LITE_MICRO_EXPECT_EQ(1, input->dims->data[0]);
TF_LITE_MICRO_EXPECT_EQ(1, input->dims->data[1]);
// The input is a 32 bit floating point value
TF_LITE_MICRO_EXPECT_EQ(kTfLiteFloat32, input->type);
```



Run Inference and Get output

```
// Provide an input value  
input->data.f[0] = 0.;
```

```
// Run the model on this input and check that it succeeds  
TfLiteStatus invoke_status = interpreter.Invoke();  
TF_LITE_MICRO_EXPECT_EQ(kTfLiteOk, invoke_status);
```

```
// Obtain a pointer to the output tensor  
TfLiteTensor* output = interpreter.output(0);
```

```
// Obtain the output value from the tensor  
float value = output->data.f[0];  
// Check that the output value is within 0.05 of  
// the expected value  
TF_LITE_MICRO_EXPECT_NEAR(0., value, 0.05);
```